

MACHINA

TRINITY

The agent runtime that assumes the LLM will fail.

C++20

Safety Core

9

Defense Layers

40+

Built-in Tools

v6.5

Current Release

The Problem

Every agent framework gives an LLM a knife and hopes for the best.

-  **LLM hallucinates rm -rf /**

No rollback. State corrupted forever.

-  **System breaks at 3 AM**

No audit trail. grep through unstructured logs.

-  **Tool spawns runaway process**

No resource limits. No isolation.

-  **External LLM API goes down**

Entire agent freezes. No fallback.

These aren't edge cases. They're Tuesday.

Trinity Architecture

Three concerns. They never mix.



BODY

Execution

- Transactional Tx/Rollback
- Tool Registry + Sandbox
- Plugin hash verification
- Permission leases



DRIVER

Decision

- Heuristic selector (always)
- LLM policy (optional)
- Circuit breaker fallback
- 3-tier intent resolution



MEMORY

Record

- SHA-256 hash-chained audit
- WAL + checkpoint/recovery
- Deterministic replay
- BM25 + vector hybrid search

9 Layers of Defense

Defense-in-depth. Every layer independent. No single point of failure.

1 Tx + Rollback

State integrity

2 Hash-chained Audit

Tamper-evident history

3 Allowlists

Command restriction

4 seccomp-BPF

Kernel syscall filtering

5 Permission Leases

Single-use tokens

6

Plugin Hash Pinning

SHA-256 before dlopen

7

Capability Gates

Bitmask permissions

8

SSRF Defense

DNS rebinding prevention

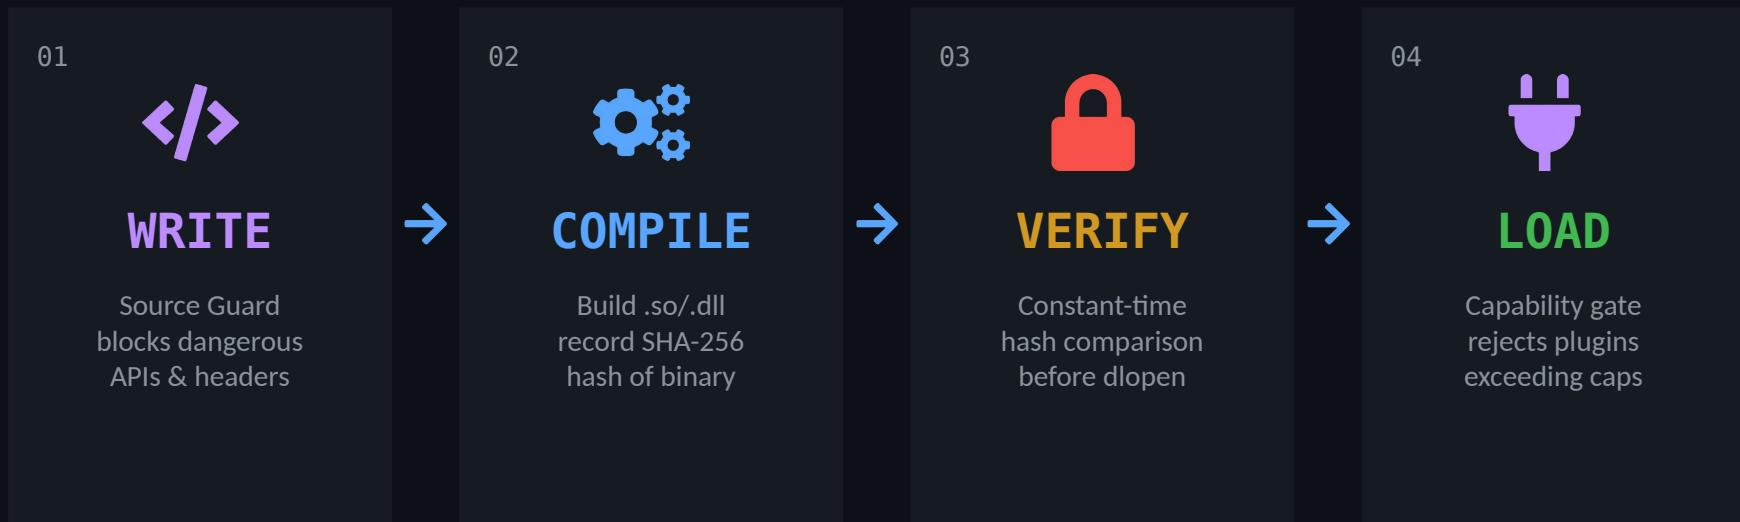
9

CRC32 WAL Framing

Crash integrity detection

Genesis: Self-Evolution

Write → Compile → Verify → Load. New tools at runtime, safely.



Opt-in only (`MACHINA_GENESIS_ENABLE=1`). Off by default in production. Three independent safety gates.

Autonomic Engine

6-level GVU cycle. The system improves itself — monotonically.



Regression Gate — bad changes blocked



Auto-Rollback — degradation auto-reverted



Monotonic — only improves, never degrades

Dual-Layer Design

C++ handles safety. Python handles intelligence.

C++ Safety Core

- Transactional execution (Tx/Rollback)
- Hash-chained audit logs (SHA-256)
- seccomp-BPF syscall filtering
- Plugin system (.so/.dll + hash pinning)
- WAL + checkpoint recovery
- Concurrent Priority Queue
- HMAC auth + rate limiting
- Prometheus /metrics endpoint

Python Agent Runtime

- Telegram bot (Pulse Loop)
- 6-level Autonomic Engine
- ExpeL + Reflexion + Distillation
- Graph Memory 2.0 (multi-hop BFS)
- MCP bridge (external tools)
- 70+ tool aliases (Korean/English)
- 3-tier intent resolution
- 36 files, all \leq 620 lines

How It Compares

Out-of-the-box capabilities.

Feature	Machina	LangChain	AutoGPT	CrewAI
Transactional execution	✓	—	—	—
Cryptographic audit trail	✓	—	—	—
Deterministic replay	✓	—	—	—
Kernel-level sandboxing	✓	—	—	—
Permission leases	✓	—	—	—
Plugin hash verification	✓	—	—	—
Circuit breaker	✓	—	—	—
Runtime self-evolution	✓	—	—	—
Prometheus /metrics	✓	—	—	—
Native C++ performance	✓	—	—	—

Test Coverage

14

C++ Unit
Test Suites

34

Python E2E
Test Cases

10

Golden
Replay Tests

39

Simulation
Scenarios

C++ Suites

CPQ · WAL · WAL Rotation · Tx · Tx Patch · Memory · Memory Query ·
Toolhost · GoalRegistry · Input Safety · Sandbox · Lease · Config · Plugin Hash

E2E Groups (13)

Chat Intent · Shell · Web Search · Code Exec · Memory · File Ops · Config ·
URL Fetch · Utilities · Chat Response · Summary · Continue Loop · Auto-
Memory

MACHINA TRINITY

Built for a world where LLMs are powerful but imperfect.

Language

C++20 + Python

License

Apache 2.0

Release

v6.5

Tests

97 total

github.com/machina-trinity